

## Exercise sheet 9

SoSe2021

Prof. Dr. Holger Fröhlich, Mohamed Aborageh, Vinay Bharadhwaj, Yasamin Salimi

**Due date: June 29th**

### Questions

#### Exercise 1 - Basics of NN (9 points)

From the MNIST database load the handwritten digits dataset.

1. Normalize your dataset before training your model. **(1 point)**
2. Train a neural network once using Adam and once using AdaGrad optimizer.  
**Hint:** Set epochs = 20, neurons of hidden layer = 100, activation function = ReLU for reproducibility. **(2 points)**
3. Plot the SparseCategoricalCrossentropy loss for both models. Plot the computed accuracy for both models. Which model performed better while training?  
**(2 points)**
4. Compute the model accuracy on the test set for both optimizers. Which model performed better? **(1 point)**
5. Familiarize yourself with Layer Normalization and explain how it works. **(1 point)**
6. Using the same dataset to train a neural network with Layer Normalization. **Hint:** Set epochs = 20, neurons of hidden layer = 100, activation function = ReLU for reproducibility.
  - a. Compute the SparseCategoricalCrossentropy loss and model accuracy.  
**(1 point)**
  - b. Evaluate the model performance using the test dataset. **(1 point)**

#### Exercise 2 - Hyper Parameter Optimization (9 points)

1. What are the main challenges with hyper-parameter optimization for neural networks? **(1 point)**
2. Inform yourself about variants of Bayesian-HPO and explain them in detail  
**(2 points)**
3. Using the same MNIST dataset, optimize the activation function for the output layer and the number of dropout units in the NN model using the following methods. **(6 points)**
  - a. Grid search
  - b. Random search
  - c. Bayesian Hyper-parameter optimization

### Exercise 3 - Transfer Learning & CNNs (7 points)

1. Load the VGG16 pre-trained model using Keras Applications API. Use the model to classify the dog images in *canines.zip* after pre-processing each image by doing the following: **(2 points)**
  - a. Load each image and set the size to 224 x 224 pixels
  - b. Convert the image pixels to a numpy array and reshape it according to the model's input requirements
  - c. Use the model to print out the predicted class and its probability for each image
2. Downscale the given matrix by applying the following pooling operations:
  - a. Max Pool **(1 point)**
  - b. Average Pool **(1 point)**

1	4	1	5
4	9	4	8
4	5	4	3
6	5	7	4

3. Load the CIFAR10 dataset using Keras datasets API and normalize the images' pixel values. Train a convolutional neural network to classify the dataset images with the following architecture: **(3 points)**
  - a. Convolutional Base:
    - i. An input convolution layer with 32 filters and a kernel size of (3,3). Adjust your input shape to that of the CIFAR images' format
    - ii. 2 convolution layers, each with 64 filters and a kernel size of (3,3)
    - iii. 2 Max Pool layers, with a pool size of 2x2
  - b. 2 dense layers, with 64 and 10 units respectively. Adjust the output of the convolutional base such that it satisfies the input requirements of the dense layers.
  - c. Use the following parameters to train the network:
    - i. Sparse categorical cross entropy as your loss function
    - ii. Adam optimizer
    - iii. 10 epochs
    - iv. ReLU activation for your layers

Compile your model, then plot the accuracy across each epoch.